

Computer Networks Assignment 3

Souparno Ghose (2022506)

October 26, 2024

Question 1

Set up four VMs as shown in the figure. Use the same setup for the entire assignment

- (a) Configure the IP addresses and routes for all VMs, as shown in the figure
- (b) Configure VM2 as the gateway such that it can forward the incoming traffic to one of the servers – add forwarding functionality

Solution

- (a) Four virtual machines (VMs) were created in VirtualBox, each configured with 1 CPU core, 512 MB of RAM, a 20 GB hard disk, and Debian (64-bit) as the operating system. The VMs were named *client*, *gateway*, *server1*, and *server2*.

I set up two host-only networks within VirtualBox. The first network was assigned the IP range `20.1.1.0/24`, and the second network was assigned `40.1.1.0/24`, both with DHCP disabled. These networks were labeled *VirtualBox Host-Only Ethernet Adapter* and *VirtualBox Host-Only Ethernet Adapter #2*, respectively.

To configure the network connections for the VMs using Netplan, I first downloaded the necessary packages for Netplan. All VMs were then set to use host-only adapters. Using the command `ip a`, I identified the available network interfaces on each VM.

I created and edited the configuration file for Netplan by running:

```
sudo nano /etc/netplan/00-installer-config.yaml
```

In the YAML file, I specified the network configurations for each VM. For instance, the *client* VM was connected to the *VirtualBox Host-Only Ethernet Adapter*, while the *gateway* VM had two interfaces: one for each of the host-only networks. The *server1* and *server2* VMs were both connected to the second host-only network.

After setting up the YAML configuration, I started the systemd network service with the following command:

```
sudo systemctl start systemd-networkd
```

I then enabled the service to start on boot:

```
sudo systemctl enable systemd-networkd
```

To apply the network configuration, I executed:

```
sudo netplan apply
```

Finally, I restarted the networking service on all VMs to ensure the changes took effect:

```
sudo systemctl restart networking.service
```

By disabling DHCP on both host-only networks and manually configuring the IP addresses, I ensured that the IP addresses remained static and consistent across reboots. This setup is crucial for networking experiments and custom routing configurations where predictable IP addresses are essential. The use of Netplan allowed for a streamlined and modern approach to managing network configurations in the Debian environment.

The network interface names in this setup, such as `enp0s3` and `enp0s8`, adhere to the predictable network interface naming convention in Linux, ensuring consistency across reboots and hardware configurations. In this naming convention:

- `en` denotes an Ethernet interface.
- `p0` indicates the physical slot number on the system's motherboard where the network device is connected.
- `s3` or `s8` represents the specific number or location of the interface on the device.

This naming scheme replaces the older `eth0`, `eth1`, etc., to enhance predictability, especially when multiple network interfaces are involved.




Figure 1: Client Configuration




Figure 2: Gateway Configuration




Figure 3: Server 1 Configuration




Figure 4: Server 2 Configuration

- (b) Initially, when attempting to ping the IP addresses 40.1.1.1 (server1) and 40.1.1.3 (server2) from the *client* VM, the pings failed. This was because the *gateway* VM (VM2) had not been configured to forward traffic between the two networks it is connected to.

To enable packet forwarding on the *gateway* VM, the following command was executed:

```
/sbin/sysctl -w net.ipv4.ip_forward=1
```

This command temporarily enables IPv4 forwarding on the system, allowing the *gateway* to forward incoming traffic from the 20.1.1.0/24 network to the 40.1.1.0/24 network, and vice versa. After executing this command, both 40.1.1.1 and 40.1.1.3 became reachable from the *client* VM, as the *gateway* was now able to route the packets properly between the two networks.




Figure 5: Showing how ping in not working before forwarding is set




Figure 6: After forwarding is set

Question 2

Traffic filtering at the gateway VM

- The gateway must block all traffic (except for ping) destined to the server 40.1.1.1/24. Show that this works; attach the screenshot.
- The gateway must block only TCP traffic initiated by 20.1.1.1/24. Show that this works; attach the screenshot.

Solution

(a) Blocking all traffic except for ping to the server 40.1.1.1/24

To accomplish the task of allowing only ping (ICMP) traffic to server 40.1.1.1 while blocking all other types of traffic, the following iptables rules were applied on the *gateway* VM:

```
/sbin/iptables -A FORWARD -d 40.1.1.1 -p icmp -j ACCEPT
/sbin/iptables -A FORWARD -d 40.1.1.1 -j DROP
```

The first rule explicitly allows ICMP (ping) traffic destined for 40.1.1.1. The second rule blocks all other traffic to 40.1.1.1, regardless of protocol, by dropping the packets. This ensures that tools like `curl`, `wget`, or any non-ICMP traffic (such as TCP or UDP) will be blocked, while pings (ICMP echo requests) will still succeed.

After applying these rules, I was able to successfully ping 40.1.1.1 from the *client* VM, but any attempt to access the server using other protocols (like HTTP or SSH) was blocked, confirming the functionality of the rule.

(b) Blocking only TCP traffic initiated by 20.1.1.1/24

To block only TCP traffic originating from the 20.1.1.1/24 network (the *client* VM's subnet), while allowing other types of traffic like UDP, the following `iptables` rule was applied on the *gateway* VM:

```
/sbin/iptables -A FORWARD -s 20.1.1.1 -p tcp -j DROP
```

This rule drops all TCP traffic originating from the *client* VM with source IP 20.1.1.1, preventing the *client* from establishing any TCP connections, such as HTTP or SSH, to any other machine. However, non-TCP traffic, such as UDP, remains unaffected and can still be forwarded.

After applying this rule, any TCP-based connections (e.g., trying to SSH or access a web server) failed, but UDP traffic (such as `traceroute` or other UDP-based protocols) still worked, demonstrating the effectiveness of the filtering rule.




Figure 7: Q2 Part A Ping is working but traceroute is not




Figure 8: Q2 Part B TCP Packets are not sent while UDP Packets are being sent

Question 3

Use the configuration obtained in Q.2. to solve this question

- Use “iperf2” tool to test the TCP and UDP bandwidth between 20.1.1.1/24 and 40.1.1.3/24. Attach the screenshot.
- What is the minimum, average, and maximum RTT (Attach the screenshot) (i) from 20.1.1.1/24 to 40.1.1.1/24 (ii) from 20.1.1.1/24 to 40.1.1.3/24 (iii) Did you find a significant difference between (i) and (ii)? If so, why?

Solution

(a) Testing TCP and UDP Bandwidth between 20.1.1.1/24 and 40.1.1.3/24

To test the TCP and UDP bandwidth between the *client* (20.1.1.1) and *server2* (40.1.1.3), the *iperf2* tool was used.

TCP Bandwidth Test: For the TCP bandwidth test, the following commands were executed:

- On *server2* (40.1.1.3), the server mode of *iperf2* was initiated:

```
iperf -s
```

- On the *client* (20.1.1.1), the following command was used to run the bandwidth test:


```
iperf -c 40.1.1.3
```

However, no results were obtained from this test because, as per the configuration applied in Q.2(b), all TCP packets originating from 20.1.1.1 are blocked by the following `iptables` rule:

```
/sbin/iptables -A FORWARD -s 20.1.1.1 -p tcp -j DROP
```

This rule drops all TCP traffic from the *client*, preventing any TCP connection from being established for the bandwidth test.

UDP Bandwidth Test: For the UDP bandwidth test, the following commands were used:

- On *server2* (40.1.1.3), the `iperf2` server was initiated in UDP mode:

```
iperf -u -s
```

- On the *client* (20.1.1.1), the UDP bandwidth test was initiated with the following command:

```
iperf -u -c 40.1.1.3
```

The UDP test showed a successful data transfer with a measured bandwidth of approximately **1.05 Mbits/sec**. Since UDP traffic is not affected by the TCP blocking rule set in Q.2(b), the packets were able to flow between the client and server without restriction.

(b) Measuring Round Trip Time (RTT)

To measure the minimum, average, and maximum Round Trip Time (RTT), `ping` was used from the *client* (20.1.1.1) to both *server1* (40.1.1.1) and *server2* (40.1.1.3).

- (i) From 20.1.1.1 to 40.1.1.1, the minimum, average, and maximum RTT were measured with `ping`.
- (ii) From 20.1.1.1 to 40.1.1.3, the minimum, average, and maximum RTT were also measured using `ping`.

(iii) RTT Comparison: There was no significant difference in the minimum, average, or maximum RTT between the pings sent to 40.1.1.1 and those sent to 40.1.1.3. The similarity in RTTs can be attributed to the fact that both servers are connected to the same network segment 40.1.1.0/24, resulting in similar physical distances and network latencies. The gateway handles traffic between the 20.1.1.0/24 and 40.1.1.0/24 networks equally for both servers, hence there is no significant difference in the RTT measurements.




Figure 9: Q3 Part A (TCP Bandwidth)

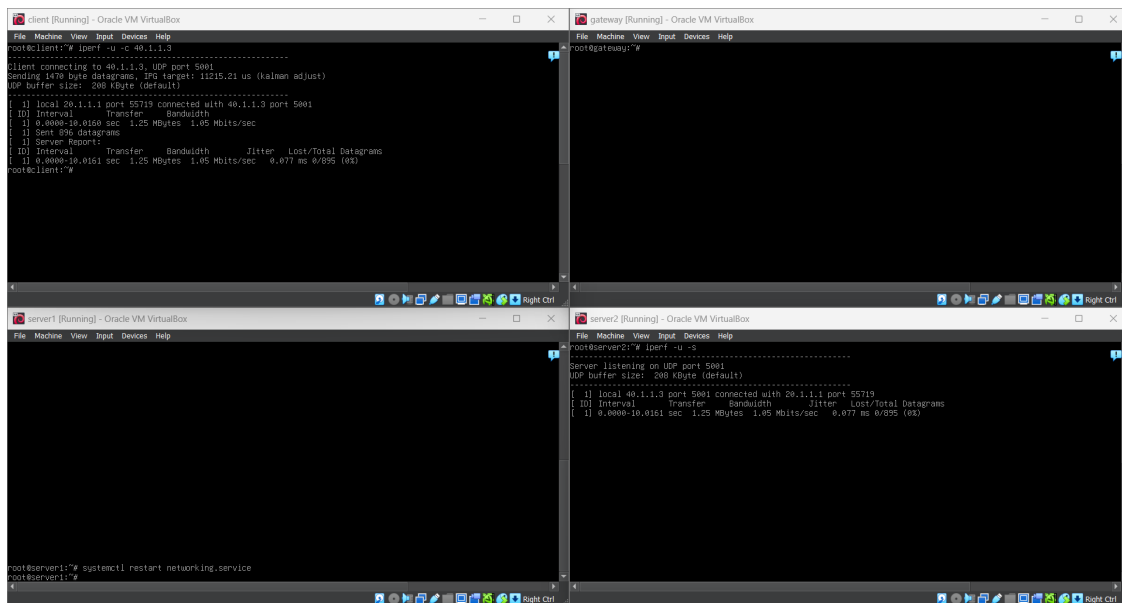


Figure 10: Q3 Part A (UDP Bandwidth)

```
root@client:~# ping -c 10 40.1.1.1
PING 40.1.1.1 (40.1.1.1) 56(84) bytes of data.
64 bytes from 40.1.1.1: icmp_seq=1 ttl=63 time=92.4 ms
64 bytes from 40.1.1.1: icmp_seq=2 ttl=63 time=1.17 ms
64 bytes from 40.1.1.1: icmp_seq=3 ttl=63 time=0.985 ms
64 bytes from 40.1.1.1: icmp_seq=4 ttl=63 time=1.39 ms
64 bytes from 40.1.1.1: icmp_seq=5 ttl=63 time=1.23 ms
64 bytes from 40.1.1.1: icmp_seq=6 ttl=63 time=1.67 ms
64 bytes from 40.1.1.1: icmp_seq=7 ttl=63 time=1.28 ms
64 bytes from 40.1.1.1: icmp_seq=8 ttl=63 time=1.60 ms
64 bytes from 40.1.1.1: icmp_seq=9 ttl=63 time=1.16 ms
64 bytes from 40.1.1.1: icmp_seq=10 ttl=63 time=1.52 ms

--- 40.1.1.1 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9016ms
rtt min/avg/max/mdev = 0.985/10.438/92.392/27.318 ms
root@client:~#
```

Figure 11: Q3 Part B Pinging 40.1.1.1

```
root@client:~# ping -c 10 40.1.1.3
PING 40.1.1.3 (40.1.1.3) 56(84) bytes of data.
64 bytes from 40.1.1.3: icmp_seq=1 ttl=63 time=90.3 ms
64 bytes from 40.1.1.3: icmp_seq=2 ttl=63 time=1.18 ms
64 bytes from 40.1.1.3: icmp_seq=3 ttl=63 time=1.89 ms
64 bytes from 40.1.1.3: icmp_seq=4 ttl=63 time=1.48 ms
64 bytes from 40.1.1.3: icmp_seq=5 ttl=63 time=1.14 ms
64 bytes from 40.1.1.3: icmp_seq=6 ttl=63 time=1.73 ms
64 bytes from 40.1.1.3: icmp_seq=7 ttl=63 time=1.13 ms
64 bytes from 40.1.1.3: icmp_seq=8 ttl=63 time=1.20 ms
64 bytes from 40.1.1.3: icmp_seq=9 ttl=63 time=1.82 ms
64 bytes from 40.1.1.3: icmp_seq=10 ttl=63 time=1.17 ms

--- 40.1.1.3 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9013ms
rtt min/avg/max/mdev = 1.128/10.303/90.301/26.667 ms
```

Figure 12: Q3 Part B Pinging 40.1.1.3

Question 4

Network address translation at the gateway VM

- (a) Change the source IP address of every packet from 20.1.1.1/24 to 40.1.1.2/24
- (b) When the packet response for the packet from step “a” arrives at the gateway, revert the destination IP address to the original.
- (c) Validate the above by sending traffic and observing the packets at each VM using Wireshark/tcpdump. Attach the screenshot.

Solution

Flushing Existing iptables Rules

Before configuring the Network Address Translation (NAT) rules, the existing `iptables` rules were flushed using the following command:

```
/sbin/iptables -F
```

This ensures that the system starts with a clean state, free of any previously set firewall or NAT rules.

(a) Changing the Source IP Address

To change the source IP address of every outgoing packet from the 20.1.1.1/24 network to 40.1.1.2/24 (the *gateway* VM’s IP), the following `iptables` rule was applied:

```
/sbin/iptables -t nat -A POSTROUTING -s 20.1.1.1 -j SNAT --to-source 40.1.1.2
```

This rule is part of the NAT table and applies source network address translation (SNAT) to packets originating from 20.1.1.1. It changes the source IP address to 40.1.1.2 when the packet is routed out of the gateway. This allows the server receiving the packet to believe it came from 40.1.1.2 instead of the *client* VM’s real IP.

(b) Reverting the Destination IP Address on Response

When a response packet destined for 40.1.1.2 (modified in step (a)) returns to the gateway, we need to revert the destination IP address to the original source, 20.1.1.1. This is accomplished using the following `iptables` rule:

```
/sbin/iptables -t nat -A PREROUTING -d 40.1.1.2 -j DNAT --to-destination 20.1.1.1
```

This rule performs destination network address translation (DNAT) in the `PREROUTING` chain. It changes the destination IP address from 40.1.1.2 back to 20.1.1.1 so that the response reaches the original client VM.

(c) Validating the NAT Configuration with `tcpdump` and `netcat`

To validate the above NAT configuration, traffic was sent between the `client` (20.1.1.1) and `server1` (40.1.1.1) while observing the packets on both systems using `tcpdump`. The following steps were taken:

- On 40.1.1.1, I started capturing the TCP packets using:

```
tcpdump tcp &
```

- On 20.1.1.1, I also initiated a packet capture:

```
tcpdump tcp &
```

- Next, I used `netcat` (`nc`) to create a simple TCP connection. On 40.1.1.1, I started a listener on port 8080:

```
nc -l -p 8080
```

- On 20.1.1.1, I initiated a connection to 40.1.1.1 using `nc`:

```
nc 40.1.1.1 8080
```

- Using `tcpdump`, I observed the routes taken by the packets. The source IP was initially modified from 20.1.1.1 to 40.1.1.2 when the packets reached the gateway. On their return, the destination IP was reverted back to 20.1.1.1, confirming that the NAT rules were correctly applied and the communication was valid.

Question 5

Load balancing at the gateway VM. Attach screenshots

- (a) Using the information obtained from Q.5.b., balance the traffic from 20.1.1.1/24 to the servers, 40.1.1.1/24 and 40.1.1.3/24. The probability of assigning the packet to the servers is 0.8 and 0.2, i.e., assign a high probability to the server with lower RTT.
- (b) Test the above configuration using a series of “ping” packets.

Solution

Before setting up the load balancing rules, I flushed the existing NAT table from the configuration of Q.4 using the following command:

```
/sbin/iptables -t nat -F
```

This ensures that no previous rules interfere with the new load balancing configuration.

(a) Configuring Load Balancing with 80-20 Traffic Split

To balance traffic from the `client` (20.1.1.1/24) to `server1` (40.1.1.1/24) and `server2` (40.1.1.3/24) based on the RTT, I used the following `iptables` rules:

```
/sbin/iptables -A PREROUTING -t nat -d 20.1.1.2 -m statistic --mode random --probability 0.2 -j DNAT --to-destination 40.1.1.1
/sbin/iptables -A PREROUTING -t nat -d 20.1.1.2 -j DNAT --to-destination 40.1.1.3
```

These two rules perform Destination Network Address Translation (DNAT) in the `PREROUTING` chain. The first rule sends 20% of the traffic to `server1` (40.1.1.1) by using the `--probability 0.2` option, and the second rule sends the remaining 80% of the traffic to `server2` (40.1.1.3), which had a lower RTT based on prior measurements.

To ensure that the response packets are returned with the correct source IP address, the following `iptables` rules were applied to perform Source Network Address Translation (SNAT) in the `POSTROUTING` chain:

```
/sbin/iptables -A POSTROUTING -t nat -d 40.1.1.1 -j SNAT --to-source 40.1.1.2
/sbin/iptables -A POSTROUTING -t nat -d 40.1.1.3 -j SNAT --to-source 40.1.1.2
```

These rules modify the source IP address of the outgoing packets to 40.1.1.2 so that the responses are directed back to the gateway VM.

Through the use of `ping` testing, it was observed that 40.1.1.3 was less congested with lower RTT values compared to 40.1.1.1, confirming that 80% of the traffic was directed to `server2` (40.1.1.3), as expected.

(b) Validating Load Balancing using ping and tcpdump

To validate the traffic distribution, I used the following commands to capture ICMP traffic on both servers:

- On **server1** (40.1.1.1), I ran:

```
tcpdump icmp &
```

- On **server2** (40.1.1.3), I ran:

```
tcpdump icmp &
```

On the *client* (20.1.1.1), I used the following command to send 10 ping requests to 20.1.1.2:

```
ping -c 1 20.1.1.2
```

After executing the command 10 times, I observed that out of 10 ICMP echo requests, 8 packets were sent to **server2** (40.1.1.3) and 2 packets were sent to **server1** (40.1.1.1), confirming the expected 80-20 traffic split based on RTT.




Figure 15: Q5 Part A

The figure displays four terminal windows from Oracle VM VirtualBox, arranged in a 2x2 grid. The top-left window, titled 'client [Running]', shows the execution of a series of ping commands from a client to a server at IP 20.1.1.2. The output includes statistics for each ping, such as '1 packets transmitted, 1 received, 0% packet loss, time 0ms' and 'rtt min/avg/max/mdev = 2.359/2.359/2.359/0.000 ms'. The top-right window, titled 'gateway [Running]', shows the configuration of a gateway. It includes commands to set up IP tables for NAT, prerouting, postrouting, and masquerading, and a command to echo the configuration to a file. The bottom-left window, titled 'server1 [Running]', shows the execution of a tcpdump command to capture ICMP traffic on the server. The output shows a list of captured packets, including echo requests and replies, with details like 'listening on enp0s3, link-type EN10MB (Ethernet), snapshot length 262144 bytes'. The bottom-right window, titled 'server2 [Running]', shows the execution of a tcpdump command to capture ICMP traffic on the server. The output shows a list of captured packets, including echo requests and replies, with details like 'listening on enp0s3, link-type EN10MB (Ethernet), snapshot length 262144 bytes'.

Figure 16: Q5 Part B